

Predicting heart disease using machine learning

Objective

Given details of patients, can we predict whether or not they have heart disease?

Data

Data is collected from : <https://archive.ics.uci.edu/ml/datasets/heart+Disease>

Features

There are following features on our data:

Create data dictionary

1. **age** - age in years
2. **sex** - (1 = male; 0 = female)
3. **cp** - chest pain type
 - 0: Typical angina: chest pain related decrease blood supply to the heart
 - 1: Atypical angina: chest pain not related to heart
 - 2: Non-anginal pain: typically esophageal spasms (non heart related)
 - 3: Asymptomatic: chest pain not showing signs of disease
4. **trestbps** - resting blood pressure (in mm Hg on admission to the hospital) anything above 130-140 is typically cause for concern
5. **chol** - serum cholestoral in mg/dl
 - serum = LDL + HDL + .2 * triglycerides
 - above 200 is cause for concern
6. **fbs** - (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)
 - '>126' mg/dL signals diabetes
7. **restecg** - resting electrocardiographic results
 - 0: Nothing to note
 - 1: ST-T Wave abnormality
 - can range from mild symptoms to severe problems
 - signals non-normal heart beat
 - 2: Possible or definite left ventricular hypertrophy
 - Enlarged heart's main pumping chamber
8. **thalach** - maximum heart rate achieved
9. **exang** - exercise induced angina (1 = yes; 0 = no)
10. **oldpeak** - ST depression induced by exercise relative to rest looks at stress of heart during exercise unhealthy heart will stress more
11. **slope** - the slope of the peak exercise ST segment
 - 0: Upsloping: better heart rate with exercise (uncommon)

- 1: Flatsloping: minimal change (typical healthy heart)
 - 2: Downsloping: signs of unhealthy heart
12. **ca** - number of major vessels (0-3) colored by fluoroscopy
- colored vessel means the doctor can see the blood passing through
 - the more blood movement the better (no clots)
13. **thal** - thalium stress result
- 1,3: normal
 - 6: fixed defect: used to be defect but ok now
 - 7: reversible defect: no proper blood movement when exercising
14. **target** - have disease or not (1=yes, 0=no) (= the predicted attribute)

Import packages

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import train_test_split, cross_val_score,
RandomizedSearchCV, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import RocCurveDisplay, ConfusionMatrixDisplay
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score
```

Import Data

```
file_path = '/kaggle/input/heart-disease-dataset/heart.csv'
df = pd.read_csv(file_path)
df.head()
```

age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
52	1	0	125	212	0	1	168	0	1.0
53	1	0	140	203	1	0	155	1	3.1
70	1	0	145	174	0	1	125	1	2.6
61	1	0	148	203	0	1	161	0	0.0
62	0	0	138	294	1	1	106	0	1.9

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

Exploratory Data Analysis

```
df.shape
```

```
(1025, 14)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

Are there any missing values?

```
df.isna().sum()
```

age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0

```
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

Inference : No missing value

```
df.describe()
```

	age	sex	cp	trestbps	chol
\count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	fbs	restecg	thalach	exang	oldpeak
\count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	0.149268	0.529756	149.114146	0.336585	1.071512
std	0.356527	0.527878	23.005724	0.472772	1.175053
min	0.000000	0.000000	71.000000	0.000000	0.000000
25%	0.000000	0.000000	132.000000	0.000000	0.000000
50%	0.000000	1.000000	152.000000	0.000000	0.800000
75%	0.000000	1.000000	166.000000	1.000000	1.800000
max	1.000000	2.000000	202.000000	1.000000	6.200000

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000

mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

Count of having heart disease or not


```
df['target'].value_counts()

target
1    526
0    499
Name: count, dtype: int64

fig, ax = plt.subplots()

x = ["Heart Disease", "No Heart Disease"]
counts = [df['target'].value_counts()[0], df['target'].value_counts()[1]]

ax.bar(x, counts, color=['salmon', 'lightblue'])
ax.set_title("Count of patients having heart disease or not")
ax.set_ylabel("Count of patients");
```



Count of Heart Disease sample = 165 count of No Heart Disease sample = 138

Inference : Dataset is balanced as number of samples of each class are roughly equal.

Heart disease counts according to sex

```
pd.crosstab(df['sex'], df['target'])
```

target	0	1
sex		
0	86	226
1	413	300


Inference : Above table shows that 75% chance of heart disease in female and 45% chance of heart disease in male.

We can deduce the same conclusion by drawing graph.

```
pd.crosstab(df['sex'], df['target']).plot(kind='bar',
color=['lightblue', 'salmon'], figsize=(10,6))

plt.title("Heart disease counts for sex")
plt.xlabel('0 = female, 1 = male')
```


```
plt.ylabel('Count')
plt.legend(['No Disease', 'Disease'])
plt.xticks(rotation=0);
```



Age vs. Max heart Rate for Heart Disease


```
import matplotlib.colors as mcolors
fig, ax = plt.subplots(figsize=(10,6))

cmap = mcolors.ListedColormap(['lightblue', 'salmon'])
scatter = ax.scatter(df['age'], df['thalach'], c=df['target'],
cmap=cmap)
ax.set_title("Heart disease in function of age and max heart rate")
ax.set_xlabel("Age")
ax.set_ylabel("Max Heart Rate")
ax.legend(*scatter.legend_elements(), title="Target");
```




Inference : In younger age groups, heart disease is more commonly found in individuals with higher heart rates, while in older age groups, heart disease is prevalent across all heart rates.

```
# Distribution of age using histogram.  
df['age'].plot(kind='hist');
```

Correlation

```
plt.figure(figsize=(12,8))  
sns.heatmap(df.corr(numeric_only=True), annot=True, fmt='.2f',  
cmap="YlGnBu");
```



Inference : There is no such strong correlation.

Modelling

```
df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
slope \										
0	52	1	0	125	212	0	1	168	0	1.0
2										
1	53	1	0	140	203	1	0	155	1	3.1
0										
2	70	1	0	145	174	0	1	125	1	2.6
0										
3	61	1	0	148	203	0	1	161	0	0.0
2										
4	62	0	0	138	294	1	1	106	0	1.9
1										
	ca	thal	target							
0	2	3	0							
1	0	3	0							
2	0	3	0							

```
3    1    3    0
4    3    2    0
```

```
# Split the data into X and y
X = df.drop('target', axis = 1)
y = df['target']
```

X

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
oldpeak \									
0	52	1	0	125	212	0	1	168	0
1.0									
1	53	1	0	140	203	1	0	155	1
3.1									
2	70	1	0	145	174	0	1	125	1
2.6									
3	61	1	0	148	203	0	1	161	0
0.0									
4	62	0	0	138	294	1	1	106	0
1.9									
...
...									
1020	59	1	1	140	221	0	1	164	1
0.0									
1021	60	1	0	125	258	0	0	141	1
2.8									
1022	47	1	0	110	275	0	0	118	1
1.0									
1023	50	0	0	110	254	0	0	159	0
0.0									
1024	54	1	0	120	188	0	1	113	0
1.4									

	slope	ca	thal
0	2	2	3
1	0	0	3
2	0	0	3
3	2	1	3
4	1	3	2
...
1020	2	0	2
1021	1	1	3
1022	1	1	2
1023	2	0	2
1024	1	1	3

[1025 rows x 13 columns]

y

```

0      0
1      0
2      0
3      0
4      0
..
1020   1
1021   0
1022   0
1023   1
1024   0
Name: target, Length: 1025, dtype: int64

```

```

np.random.seed(42)

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2)

X_train

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
oldpeak \									
835	49	1	2	118	149	0	0	126	0
0.8									
137	64	0	0	180	325	0	1	154	1
0.0									
534	54	0	2	108	267	0	0	167	0
0.0									
495	59	1	0	135	234	0	1	161	0
0.5									
244	51	1	2	125	245	1	0	166	0
2.4									
..
..									
700	41	1	2	130	214	0	0	168	0
2.0									
71	61	1	0	140	207	0	0	138	1
1.9									
106	51	1	0	140	299	0	1	173	1
1.6									
270	43	1	0	110	211	0	1	161	0
0.0									
860	52	1	0	112	230	0	1	160	0
0.0									
	slope	ca	thal						
835	2	3	2						
137	2	0	2						
534	2	0	2						
495	1	0	3						
244	1	0	2						

..
700	1	0	2
71	2	1	3
106	2	0	3
270	2	0	3
860	2	1	2

[820 rows x 13 columns]

As we know this problem is of classification type. Hence we will use these three models:

- Logistic Regression
- K-Nearest Neighbours Classifier
- Random Forest Classifier

```
def fit_and_score(models, X_train, X_test, y_train, y_test):
    """
    Fit and train models simultaneously
    """
    models_score = {}

    for key, values in models.items():
        model = values
        model.fit(X_train, y_train)
        models_score[key] = model.score(X_test, y_test)

    return models_score

models = {
    "LogisticRegression": LogisticRegression(max_iter=1000),
    "KNN": KNeighborsClassifier(),
    "RandomForestClassifier": RandomForestClassifier()
}

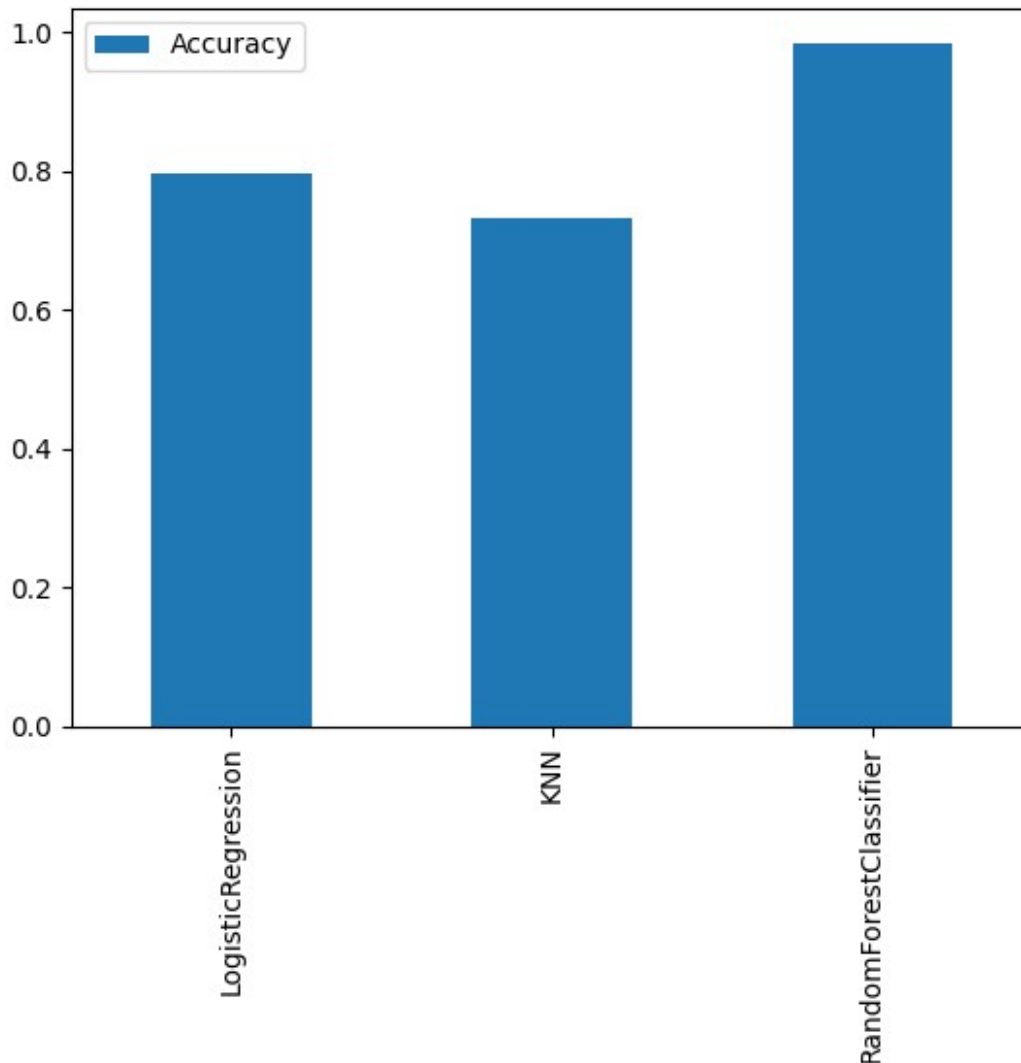
models_score = fit_and_score(models, X_train, X_test, y_train, y_test)
models_score

{'LogisticRegression': 0.7951219512195122,
 'KNN': 0.7317073170731707,
 'RandomForestClassifier': 0.9853658536585366}
```

Model Comparision

```
model_compare = pd.DataFrame(models_score, index=["Accuracy"])
model_compare.T.plot(kind='bar')

<Axes: >
```



Inference : Here we can see that **LogisticRegression** and **RandomForestClassifier** are doing well as compared to **KNN**.

Hence, we will move forward with **LogisticRegression** and **RandomForestClassifier** models. First create a function to store the different score metrics for different model for future use.

```
def pred_metrics(model, X_test, y_true):  
    # accuracy = cross_val_score(model, X, y,  
    scoring='accuracy').mean()  
    # precision = cross_val_score(model, X, y,  
    scoring='precision').mean()  
    # recall = cross_val_score(model, X, y, scoring='recall').mean()  
    # f1 = cross_val_score(model, X, y, scoring='f1').mean()  
    y_preds = model.predict(X_test)  
    accuracy = accuracy_score(y_true, y_preds)  
    precision = precision_score(y_true, y_preds)
```

```

recall = recall_score(y_true, y_preds)
f1 = f1_score(y_true, y_preds)

metric_dict = {
    "Accuracy": round(accuracy,8),
    "Precision": round(precision,8),
    "Recall": round(recall,8),
    "F1": round(f1,8),
}

return metric_dict

```

Base Logistic Regression performance

```

lr_base_model_metrics = pred_metrics(models['LogisticRegression'],
X_test, y_test)
lr_base_model_metrics

{'Accuracy': 0.79512195,
 'Precision': 0.75630252,
 'Recall': 0.87378641,
 'F1': 0.81081081}

```

Base Random Forest Classifier performance

```

rfc_base_model_metrics =
pred_metrics(models['RandomForestClassifier'], X_test, y_test)
rfc_base_model_metrics

{'Accuracy': 0.98536585,
 'Precision': 1.0,
 'Recall': 0.97087379,
 'F1': 0.98522167}

```

Hyperparameter Tuning

Hyperparameter Tuning using RandomizedSearchCV

```

# Hyperparameter grid for LogisticRegression
lr_rs_grid = {
    "solver" : ["liblinear", 'lbfgs', 'newton-cg'],
    "C" : np.logspace(-4, 4, num=20, base=10),
    "max_iter" : [1000, 1500, 2000]
}

# Hyperparameter grid for RandomForestClassifier
rfc_rs_grid = {
    "n_estimators": np.arange(10, 1000, 50),
    "max_depth": [None, 3, 5, 10],

```

```

    "min_samples_split": np.arange(2, 20, 2),
    "min_samples_leaf": np.arange(1, 20, 2)
}

```

Logistic Regression

```

lr_rs_model = RandomizedSearchCV(models['LogisticRegression'],
lr_rs_grid, cv=5, n_iter=20)
lr_rs_model.fit(X_train, y_train)

/opt/conda/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(
RandomizedSearchCV(cv=5, estimator=LogisticRegression(max_iter=1000),
n_iter=20,
                    param_distributions={'C': array([1.00000000e-04,
2.63665090e-04, 6.95192796e-04, 1.83298071e-03,
4.83293024e-03, 1.27427499e-02, 3.35981829e-02, 8.85866790e-02,
2.33572147e-01, 6.15848211e-01, 1.62377674e+00, 4.28133240e+00,
1.12883789e+01, 2.97635144e+01, 7.84759970e+01, 2.06913808e+02,
5.45559478e+02, 1.43844989e+03, 3.79269019e+03,
1.00000000e+04])),
                    'max_iter': [1000, 1500,
2000],
                    'solver': ['liblinear',
'lbfgs',
'newton-cg']})

lr_rs_model.best_params_
{'solver': 'liblinear', 'max_iter': 1500, 'C': 1.623776739188721}
lr_rs_model.score(X_test, y_test)
0.7853658536585366

lr_rs_model.best_estimator_
LogisticRegression(C=1.623776739188721, max_iter=1500,
solver='liblinear')

```



```

lr_rs_model_metrics = pred_metrics(lr_rs_model.best_estimator_,
X_test, y_test)
lr_rs_model_metrics

{'Accuracy': 0.78536585,
 'Precision': 0.74380165,
 'Recall': 0.87378641,
 'F1': 0.80357143}

```

Random Forest Classifier

```

rfc_rs_model = RandomizedSearchCV(models['RandomForestClassifier'],
rfc_rs_grid, cv=5, n_iter=20)
rfc_rs_model.fit(X_train, y_train)

RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(),
n_iter=20,
                    param_distributions={'max_depth': [None, 3, 5, 10],
                    'min_samples_leaf': array([ 1,
3, 5, 7, 9, 11, 13, 15, 17, 19])},
                    'min_samples_split':
array([ 2, 4, 6, 8, 10, 12, 14, 16, 18])},
                    'n_estimators': array([ 10,
60, 110, 160, 210, 260, 310, 360, 410, 460, 510, 560, 610,
660, 710, 760, 810, 860, 910, 960])})

rfc_rs_model.best_params_

{'n_estimators': 960,
 'min_samples_split': 4,
 'min_samples_leaf': 3,
 'max_depth': 10}

rfc_rs_model.score(X_test, y_test)

0.9560975609756097

rfc_rs_model_metrics = pred_metrics(rfc_rs_model.best_estimator_,
X_test, y_test)
rfc_rs_model_metrics

{'Accuracy': 0.95609756,
 'Precision': 0.95192308,
 'Recall': 0.96116505,
 'F1': 0.95652174}

lr_rs_model_metrics

{'Accuracy': 0.78536585,
 'Precision': 0.74380165,


```

```
'Recall': 0.87378641,  
'F1': 0.80357143}
```

Compare all four models

(Base logisticRegression, Base RandomForestClassifier, RandomizedLogisisticRegression and RandomizedRandomForesClassifier)

```
compare_base_and_rs_metrics = pd.DataFrame({  
    "lr_base": lr_base_model_metrics,  
    "rfc_base": rfc_base_model_metrics,  
    "lr_rs": lr_rs_model_metrics,  
    "rfc_rs": rfc_rs_model_metrics  
})  
  
compare_base_and_rs_metrics  
  
   lr_base  rfc_base  lr_rs  rfc_rs  
Accuracy  0.795122  0.985366  0.785366  0.956098  
Precision  0.756303  1.000000  0.743802  0.951923  
Recall     0.873786  0.970874  0.873786  0.961165  
F1         0.810811  0.985222  0.803571  0.956522  
  
compare_base_and_rs_metrics.T.plot.bar()  
<Axes: >
```



As we can see that lr_rs(LogisticRegression from randomizedSearchCV) has more accuracy(If we emphasis more on accuracy than other metric) work better than other three.

Hence, We will move forward with LogisticRegression and apply GridSearchCV on it.

GridSearchCV

Apply GridSearchCV on LogisticRegression

```
# Hyperparameter grid for LogisticRegression(for GridSearchCV)
lr_rs_grid = {
    "solver" : ["liblinear", 'newton-cg'],
    "C" : np.logspace(-5, 5, num=30),
    "max_iter" : [1000, 1500, 2000, 2500]
}

lr_gs_model = GridSearchCV(models['LogisticRegression'], lr_rs_grid,
cv=5)
lr_gs_model.fit(X_train, y_train)

GridSearchCV(cv=5, estimator=LogisticRegression(max_iter=1000),
param_grid={'C': array([1.00000000e-05, 2.21221629e-05,
4.89390092e-05, 1.08263673e-04,
2.39502662e-04, 5.29831691e-04, 1.17210230e-03, 2.59294380e-03,
```

```
5.73615251e-03, 1.26896100e-02, 2.80721620e-02, 6.21016942e-02,
1.37382380e-01, 3.03919538e-01, 6.72335754e-01, 1.48735211e+00,
3.29034456e+00, 7.27895384e+00, 1.61026203e+01, 3.56224789e+01,
7.88046282e+01, 1.74332882e+02, 3.85662042e+02, 8.53167852e+02,
1.88739182e+03, 4.17531894e+03, 9.23670857e+03, 2.04335972e+04,
4.52035366e+04, 1.00000000e+05]),
      'max_iter': [1000, 1500, 2000, 2500],
      'solver': ['liblinear', 'newton-cg']})
```

```
lr_gs_model.best_params_
```

```
{'C': 1.4873521072935119, 'max_iter': 1000, 'solver': 'liblinear'}
```

```
lr_gs_model.score(X_test, y_test)
```

```
0.7853658536585366
```

```
lr_gs_model_metrics = pred_metrics(lr_gs_model.best_estimator_,
X_test, y_test)
```

```
lr_gs_model_metrics
```

```
{'Accuracy': 0.78536585,
'Precision': 0.74380165,
'Recall': 0.87378641,
'F1': 0.80357143}
```

```
compare_base_and_rs_metrics
```

	lr_base	rfc_base	lr_rs	rfc_rs
Accuracy	0.795122	0.985366	0.785366	0.956098
Precision	0.756303	1.000000	0.743802	0.951923
Recall	0.873786	0.970874	0.873786	0.961165
F1	0.810811	0.985222	0.803571	0.956522


After comparing logisticRegression base model, RandomizedSearch logisticRegression model and GridSearch logisticRegression we found that that is not such difference. Hence we can move forward with any of three model

Model Evaluation

- Accuracy
- Area under ROC curve
- Confusion matrix
- Classification report


```
lr_final_model = lr_rs_model.best_estimator_
lr_final_model_preds = lr_final_model.predict(X_test)
```

```
ConfusionMatrixDisplay.from_estimator(lr_final_model, X_test,
y_test,colorbar=False);
```



ROC and AUC

```
RocCurveDisplay.from_estimator(lr_final_model, X_test, y_test);
```



Classification Report

```
print(classification_report(y_test,lr_final_model_preds))
```

	precision	recall	f1-score	support
0	0.85	0.70	0.76	102
1	0.74	0.87	0.80	103
accuracy			0.79	205
macro avg	0.79	0.78	0.78	205
weighted avg	0.79	0.79	0.78	205

Evaluation metrics using cross-validation

Accuracy

```
np.random.seed(42)
lr_final_model_cv_acc = cross_val_score(lr_final_model, X, y, cv=5,
scoring='accuracy').mean()
print(f'The cross-validated accuracy is
{lr_final_model_cv_acc*100:.2f}')
```

The cross-validated accuracy is 84.78

Precision

```
np.random.seed(42)
lr_final_model_cv_precision = cross_val_score(lr_final_model, X, y,
cv=5, scoring='precision').mean()
print(f'The cross-validated precision is
{lr_final_model_cv_precision:.2f}')
```

The cross-validated precision is 0.82

Recall

```
np.random.seed(42)
lr_final_model_cv_recall = cross_val_score(lr_final_model, X, y, cv=5,
scoring='recall').mean()
print(f'The cross-validated recall is {lr_final_model_cv_recall:.2f}')
```

The cross-validated recall is 0.90

F1

```
np.random.seed(42)
lr_final_model_cv_f1 = cross_val_score(lr_final_model, X, y, cv=5,
scoring='f1').mean()
print(f'The cross-validated F1 is {lr_final_model_cv_f1:.2f}')
```

The cross-validated F1 is 0.86

AUC


```
np.random.seed(42)
lr_final_model_cv_auc = cross_val_score(lr_final_model, X, y, cv=5,
scoring='roc_auc').mean()
print(f'The cross-validated AUC score is is
{lr_final_model_cv_auc:.2f}')
```

The cross-validated AUC score is is 0.92

Visualize cross-validation score

```
metric_name = ['Accuracy', 'Precision', 'Recall', 'F1', 'Roc Auc']
metric_value = [lr_final_model_cv_acc, lr_final_model_cv_precision,
lr_final_model_cv_recall, lr_final_model_cv_f1, lr_final_model_cv_auc]

bar_container = plt.bar(metric_name, metric_value)
plt.bar_label(bar_container)
plt.title("Cross-Validated Metrics");
```



Inference:

- Accuracy = 84.8 % => It means model correctly predicted the outcome(both heart disease and no heart disease) in 84.8% of the case.
- Precision = 82.1% => It means that out of prediction of heart diseases, 82.8% of them truly have cancer.
- Recall = 92.7% => It means out of all the patients have heart disease, our model predicts 92% of them having heart disease.

Here, It is excellent at identifying patients with heart disease (92.7% recall) and fairly accurate in predicting heart disease for those who truly have it (82.1% precision).

It depends upon business task that we emphasize more on which metrics (recall or precision).

Feature Importance

```
lr_final_model.coef_  
array([[ 0.01245689, -1.68544402,  0.85198764, -0.01578096, -  
0.00828761,  
        -0.20903191,  0.32008044,  0.03461353, -0.79648686, -  
0.65056533,  
        0.56750662, -0.81568847, -1.04166391]])
```




```
df.columns
Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg',
      'thalach',
      'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
      dtype='object')

feature_dict = dict(zip(df.columns, lr_final_model.coef_[0]))
feature_df = pd.DataFrame(feature_dict, index=[0])
feature_df
```

	age	sex	cp	trestbps	chol	fbs	restecg
0	0.012457	-1.685444	0.851988	-0.015781	-0.008288	-0.209032	0.32008

	thalach	exang	oldpeak	slope	ca	thal
0	0.034614	-0.796487	-0.650565	0.567507	-0.815688	-1.041664

```
feature_df.T.plot.bar(legend=False);
```



Larger the value of coefficient, the more it contributes to decision making. If value is negative then there is negative correlation and vice-versa for positive value.

Inference:

1. `sex` feature has larger negative value implies it has strong negative correlation.

It means that value of `sex` decrease (i.e 0) there is more chance of having a heart disease. We have seen earlier that female have 75% of change of heart. Hence women have more chance of heart disease.

1. `slope` feature has larger positive value implies it has positive correlation

```
pd.crosstab(df["slope"], df["target"])
```

target	0	1
slope		
0	46	28
1	324	158
2	129	340

As we can notice from the table is that as `slope` increase chance of having heart disease also increases. (same logic goes with `cp`, `exang` and `thal` etc)

Still there is lot of thing we can do

- Try different classification model (like SVM, XGboost)
- Try more hyper parameter
- Ask for more data samples.